

Edge-Based Defense in Networks

A Modification of Bloch, Chatterjee & Dutta (2023)

Chukwudi Samuel Ogbuta

Supervised by

Prof. Sylvain Béal • Prof. Emmanuel Peterlé

7 July 2026 • Besançon

Strategic adversaries on networks

The problem.


Illegal fishing reaches places no single country owns. Vessels move along sea routes between national waters and high seas, taking what isn't theirs to take.

The adversary is strategic.

Routes are chosen, not random. The weakest gap is the one used.

The defenders are many.


Neighbouring states watch the same waters. Defense is naturally shared.



Estimated global fishing activity. Orange marks suspected illegal activity — West Africa is one of three concern zones.

What is a network?

A network is a graph $G = (V, E)$: nodes V are locations, edges E are direct connections between them.



Nodes.

Locations the attacker can reach.

Edges.

The connections between locations.

The attacker.

A single adversary at a fixed node in the network.

The defenders.

One at each target. They invest to intercept.

Two traditions, one gap

	Node Defense		Edge Defense
	Centralized	Continuous	Centralized
<i>Authors</i>	Dziubiński & Goyal Cerdeiro et al.	Bloch, Chatterjee & Dutta (2023)	Washburn & Wood Mavronicolas et al.
<i>Investment</i>	Binary	Continuous	Discrete or budgeted
<i>Defenders</i>	Single planner	Many	Single planner

The gap — *Continuous, decentralized investment — placed on the connections.*

The Bloch baseline: setup

A continuous-investment game played on the network. One attacker, one defender per target, every player optimizing.

Path survival.

The attacker reaches a target only if it slips past every defender between the attacker node and that target.

Attacker's expected payoff.

Call it U . The attacker picks a mixed strategy over paths, and U sums target values weighted by how often it uses each path and the probability of getting through.

Defender's payoff.

Each defender bears the expected loss from attacks on its own target, minus the cost of its investment.

The Bloch baseline: equilibrium

Notation: b_i = target i 's value • x_i = target i 's investment • q_i = attacker's probability of targeting i

Bloch distinguishes two kinds of target. A first target is reached directly from the attacker node; a deeper target is reached only by passing through other targets that are themselves attacked.

$$x_i = 1 - \frac{U}{b_i}, \quad q_i = \frac{1}{b_i} \left(1 - \frac{U}{b_i} \right)$$

$$x_i = 1 - \frac{b_k}{b_i}, \quad q_i = \frac{b_k}{b_i U} \left(1 - \frac{b_k}{b_i} \right)$$

First-target equilibrium.

Higher-value targets invest more and are attacked more often. U is the common payoff across all attacked targets.

Deeper-target equilibrium.

For a target i reached only through its predecessor k . The predecessor's value b_k replaces U as the relevant threshold.

What Bloch leaves on the table



Defense lives at the location.

In Bloch, every defender's investment sits at its own node. The attacker is caught *at* a target, not on the way to one.

But corridors are shared.

A shipping lane, a forest road, a fishing zone — none of these belong to one party alone. Real defense happens *on the way*, by whoever watches that stretch.

Bloch: defense at nodes



Edge-based: defense on edges

The modification, stated

Move investment from nodes to edges. Every edge has two defenders.

Bloch's baseline

$$S(\text{node } i) = 1 - x_i$$

One defender per node. Survival is whatever it fails to catch.

The modification

$$S(e) = (1 - x_{a,e})(1 - x_{b,e})$$

Two defenders per edge. The attacker escapes only if *both* sides fail.

Cost is split: each defender i pays $\sum_{e \ni i} x_{i,e}^2 / 2$, summed over its incident edges.

What carries over from Bloch

Three of the four key lemmas in Bloch's analysis depend only on the attacker's optimization. They survive the move to edges.

Equal payoff across the support

Every attacked target gives the attacker the same payoff U .

Optimal predecessor has lowest value

If a target can be reached through several support members, the attacker passes through the one with the lowest value.


Values increase along attack paths

Target values rise as the attacker moves deeper. Survival decay is offset by rising value.

Plus: equilibrium exists and is unique, under the same genericity assumption Bloch uses.

What changes (1): the multi-path factor

A target reachable through several free-highway edges is structurally weaker: each edge must be defended separately.



$m_i = 3$ free-highway paths
through different non-support nodes

$$\text{Bloch: } q_i = \frac{b_i - U}{b_i^2}$$

$$\text{Edge-based: } q_i = \frac{m_i (b_i - U)}{b_i^2}$$

The factor m_i counts the free-highway edges feeding the target.
More edges $\Rightarrow$ more attack probability $\Rightarrow$ weaker target.

What changes (2): strategic substitutes

When two defenders share an edge, one's investment changes the other's incentive.

$$x_{i,e}^* = q b_i (1 - x_{j,e})$$

$$\frac{\partial x_{i,e}^*}{\partial x_{j,e}} = -q b_i < 0$$

If my neighbor invests more, I invest less.

The slope is strictly negative whenever attacks flow through the edge. Defenders' efforts substitute for each other.

But the edge cannot stay undefended.

If the other side invests zero, my best response is strictly positive. $(0, 0)$ is never an equilibrium when attacks reach the edge.

Other structural changes

Three further differences from Bloch's framework, briefly.

Mutual coverage on edges

When a target has m_i free-highway edges, the attacker spreads attack equally across them, and the defender invests equally on each.

No Braess paradox

Adding a connection in Bloch can hurt the attacker. Here, every new free-highway edge strictly raises attacker payoff U . Braess has no room to operate.

Mixed configurations

A target can be reached by both free-highway edges *and* through a predecessor. Each edge gets its own equilibrium equation. No analog in Bloch.


Managing a shared edge: four possibilities

Two endpoints, two modelling choices. Crossing them gives four cases.

	Independent interception	Joint interception
	$S(e) = (1 - x_{i,e})(1 - x_{j,e})$	$S(e) = 1 - x_{i,e} x_{j,e}$
Separable cost $x_{i,e}^2 / 2$ each	Case A (this chapter) Substitutes. Edge never undefended.	Case C Complements. $(0, 0)$ possible.
Shared cost $(x_i + x_j)^2 / 4$ each	Case B Stronger substitutes. Edge never undefended.	Case D Threshold at $q b_i = 1/2$.

Headline: technology matters more than cost — it decides whether an undefended edge can persist.

Example: a target with two paths



Values: $b_1=0.20$, $b_2=0.25$, $b_3=0.50$, $b_4=0.70$

T_4 has $m_4 = 2$ (reached via T_1 and T_2)

	Bloch	Edge-based
Attacker utility U	0.4020	0.4773 (+18.7%)
Attack on T_4 (q_4)	0.6081	0.9091
Attack on T_3 (q_3)	0.3920	0.0909
Investment x_4	0.4257	0.3181 per edge
Investment x_3	0.1960	0.0455

More edges to defend $\Rightarrow$ higher $U \Rightarrow T_3$ pushed to the margin $\Rightarrow$ attacks concentrate on T_4 .

Beyond complete information

Real players rarely observe everything. The edge-based model extends in two natural directions.

Private target values.

Each defender knows its own value b_i , drawn from a common prior F . The attacker operates under the prior alone.


Partial network observation.

Each player observes its own incident edges with certainty, and assigns probability θ to every unseen edge.

Motivated by cybersecurity: Chang et al. (2024) show attackers target firms by observable signals (size, sales) — not by the true value at stake. The two extensions map onto these two information asymmetries.

Equilibrium concept: Bayes-Nash. Worked examples: the triangle network.

Private values on the triangle



Complete info: $b_1=0.49$, $b_2=0.51$

Incomplete info: F uniform on $(0.2, 0.8)$,

$$\mu=0.5$$

	Complete	Incomplete
Attacker utility U	0.3869	0.3726
Shared edge defended by	T_2 only	both
x on $\{T_1, T_2\}$ by T_1	0.0000	0.0168
x on $\{T_1, T_2\}$ by T_2	0.0392	0.0175

Headline: uncertainty disrupts perfect coordination. Both defenders cover the shared edge — neither knows for sure they are the lower-value target. Attacker is worse off.

Partial network observation

Same triangle. Now b_1, b_2 are common knowledge but each player believes any unseen edge exists with probability θ .

θ	Regime
$\in (0, 0.961)$	Direct-only ($\gamma_1 = \gamma_2 = 0$)
$\in [0.961, 1]$	Partial roundabout ($\gamma_2 > 0$)

Threshold: $\theta^* \approx 0.961$

Below the threshold.

The shared edge isn't used by the attacker, so defenders leave it undefended. U is locked at 0.3748 — independent of how uncertain the network is.

Above the threshold.

Roundabout routes become attractive. Defender 2 starts covering the shared edge. As $\theta \rightarrow 1$, U rises smoothly to the complete-info value 0.3869.

Headline: a sharp threshold, not a smooth slope. Below it, uncertainty is “absorbing.” In larger networks, the threshold likely drops as more roundabout routes become available.

Simulator: building a network

The screenshot displays the BLOCH SIMULATOR interface, which is used for building a network. The interface is divided into two main sections: CONFIGURATION and NETWORK VISUALIZATION.

CONFIGURATION

- Network Type:** A dropdown menu set to "Line".
- Number of Targets:** A slider set to 2, with a range from 2 to 10.
- TARGET VALUES:**
 - T1 Value:** A text input field containing "0.2".
 - T2 Value:** A text input field containing "0.85".
- BUILD NETWORK:** A blue button with a lightning bolt icon.

NETWORK VISUALIZATION

The visualization shows a linear network topology with three nodes connected by a horizontal line:

- ATK:** A red circular node on the left.
- T1 0.2:** A green circular node in the middle.
- T2 0.85:** A green circular node on the right.


Build any network.

Choose topology, set targets, assign values.

Built-in checks.

Unique values, reachable targets, no orphans.

Simulator: solving the equilibrium



Two defense models.

Bloch (nodes) or edge-based.


Two solver modes.

Decentralized or cooperative.

Support reduction.

Automatic, iterative.

Simulator: running experiments



Perturb and see.

Change a value, add a connection, or remove a target.

Before and after.

Side-by-side comparison with impact summary.

Limitations

What this work does not resolve.

Single attacker only

No multi-adversary version yet. The equilibrium and cooperative analysis would both need extension.

Incomplete info on triangle only

Bayes-Nash results derived on the smallest network where the mechanism appears. Larger networks have more routes and probably more thresholds.

Interior stability open

Strategic substitutes on shared edges raise dynamic questions about how endpoints arrive at the equilibrium. Not addressed here.

Where this goes next / Key takeaways

Future directions.

- Joint-cost variation analyzed in full
- Combined value + network uncertainty
- Multi-attacker extension
- Empirical calibration: illegal fishing, logging, waste
- Network design: where should new connections go?

Key takeaways.

- Moving defense from *nodes* to *edges* produces a coherent, decentralized model — a gap that node-defense and edge-defense traditions had each left open.
- Multi-path weakness, strategic substitutes, and no Braess paradox all follow from one structural change.
- The framework survives incomplete information in two distinct settings.

Thank you.

Questions and comments very welcome.